

# Reinforcement Learning

## Lecture 5: Monte-Carlo & Temporal-Difference Learning

S. M. Ahsan Kazmi

# Recap

Last lecture:

- Planning via DP for a **known** MDP
  - Policy Evaluation
  - Policy Iteration
  - Value Iteration
  - Extensions (Dynamic Programming)

This lecture:

- Model-free prediction to estimate values in **an unknown** MDP
  - Monte-Carlo Learning
  - Temporal-Difference Learning

## Recap: Markov decision process (MDP)

- A Markov decision process (MDP) is a Markov reward process with decisions.
- It is an environment in which all states are Markov.

### Definition

A *Markov Decision Process* is a tuple  $\langle \mathcal{S}, \mathcal{A}, \mathcal{P}, \mathcal{R}, \gamma \rangle$

- $\mathcal{S}$  is a finite set of states
- $\mathcal{A}$  is a finite set of actions
- $\mathcal{P}$  is a state transition probability matrix,  
 $\mathcal{P}_{ss'}^a = \mathbb{P}[S_{t+1} = s' \mid S_t = s, A_t = a]$
- $\mathcal{R}$  is a reward function,  $\mathcal{R}_s^a = \mathbb{E}[R_{t+1} \mid S_t = s, A_t = a]$
- $\gamma$  is a discount factor  $\gamma \in [0, 1]$ .

# Monte-Carlo Learning

# Monte-Carlo Learning

- MC methods learn **directly** from **episodes of experience**
- MC is **model-free**: no knowledge of MDP transitions/rewards
- MC learns from **complete episodes**: no bootstrapping
- MC uses the simplest possible idea: **value = mean return**
  - Caveat: can only apply MC to **episodic** MDPs
- All episodes must **terminate**

# Monte-Carlo Policy Evaluation

Goal: learn  $v_\pi$  from episodes of experience under policy  $\pi$

$$S_1, A_1, R_2, \dots, S_k \sim \pi$$

The **return** is the total discounted reward (for an episode ending at time  $T > t$ ):

$$G_t = R_{t+1} + \gamma R_{t+2} + \dots + \gamma^{T-t-1} R_T$$

The value function is the expected return:

$$v_\pi(s) = \mathbb{E} [G_t \mid S_t = s, \pi]$$

We can just use **sample average** return instead of **expected** return

We call this **Monte Carlo policy evaluation**

## First-Visit Monte-Carlo Policy Evaluation

- To evaluate state  $s$
- The **first time-step**  $t$  that state  $s$  is visited in an episode
- Increment counter  $N(s) \leftarrow N(s) + 1$
- Increment total return  $S(s) \leftarrow S(s) + G_t$
- Value is estimated by mean return  $V(s) = S(s)/N(s)$
- By law of large numbers,  $V(s) \rightarrow v_\pi(s)$  as  $N(s) \rightarrow \infty$

## First-Visit Monte-Carlo Policy Evaluation

Initialize:

$\pi \leftarrow$  policy to be evaluated

$V \leftarrow$  an arbitrary state-value function

$Returns(s) \leftarrow$  an empty list, for all  $s \in \mathcal{S}$

Repeat forever:

Generate an episode using  $\pi$

For each state  $s$  appearing in the episode:

$G \leftarrow$  return following the first occurrence of  $s$

Append  $G$  to  $Returns(s)$


$V(s) \leftarrow \text{average}(Returns(s))$



## Every-Visit Monte-Carlo Policy Evaluation

- To evaluate state  $s$
- **Every** time-step  $t$  that state  $s$  is visited in an episode
- Increment counter  $N(s) \leftarrow N(s) + 1$
- Increment total return  $S(s) \leftarrow S(s) + G_t$
- Value is estimated by mean return  $V(s) = S(s)/N(s)$
- Again,  $V(s) \rightarrow v_{\pi}(s)$  as  $N(s) \rightarrow \infty$

# Example



- 4 states MC: A, B, C, and a terminal state D. Assume the policy is random. The transition reward is -1.

- Episode 1:  $A \rightarrow B \rightarrow C \rightarrow D$
- Episode 2:  $A \rightarrow C \rightarrow B \rightarrow C \rightarrow D$

## First-Visit Monte Carlo

- Episode 1:

- $V(A) = -3$
- $V(B) = -2$
- $V(C) = -1$
- $V(D) = 0$  (terminal)

- Episode 2:

- $V(A) = ((-3) + (-4)) / 2 = -3.5$
- $V(B) = ((-2) + (-2)) / 2 = -2.0$
- $V(C) = ((-1) + (-3)) / 2 = -2.0$

## Every Visit Monte Carlo

- Episode 1:

- $V(A) = -3$
- $V(B) = -2$
- $V(C) = -1$
- $V(D) = 0$  (terminal)

- Episode 2:

- $V(A) = ((-3) + (-4)) / 2 = -3.5$
- $V(B) = ((-2) + (-2)) / 2 = -2.0$
- $V(C) = (-1 + -3 + -1) / 3 = -1.67$

## Incremental Monte-Carlo Updates

- Update  $V(s)$  incrementally after episode  $S_1, A_1, R_2, \dots, S_T$
- For each state  $S_t$  with return  $G_t$

$$N(S_t) \leftarrow N(S_t) + 1$$

$$V(S_t) \leftarrow V(S_t) + \frac{1}{N(S_t)} (G_t - V(S_t))$$

- In non-stationary problems, it can be useful to track a running mean, i.e. forget old episodes.

$$V(S_t) \leftarrow V(S_t) + \alpha (G_t - V(S_t))$$

# Summary

MC has several **advantages** over DP:

- Can learn directly from interaction with the environment
- No need for full models
- No need to learn about ALL states (no bootstrapping)
- Less harmed by violating Markov property
- MC methods provide an alternate policy evaluation process

Visit **Chapter 5** for more details

# Temporal-Difference Learning

# Temporal-Difference Learning(TD)

- TD methods learn directly from episodes of experience
- TD is model-free: no knowledge of MDP transitions/rewards
- TD learns from incomplete episodes, by bootstrapping
- TD updates a **guess** towards a **guess**

## MC and TD

- Goal: learn  $v_\pi$  online from experience under policy  $\pi$
- Incremental every-visit Monte-Carlo
  - Update value  $V(S_t)$  toward *actual* return  $G_t$

$$V(S_t) \leftarrow V(S_t) + \alpha (G_t - V(S_t))$$

- Simplest temporal-difference learning algorithm: TD(0)
  - Update value  $V(S_t)$  toward *estimated* return  $R_{t+1} + \gamma V(S_{t+1})$

$$V(S_t) \leftarrow V(S_t) + \alpha (R_{t+1} + \gamma V(S_{t+1}) - V(S_t))$$

- $R_{t+1} + \gamma V(S_{t+1})$  is called the *TD target*
- $\delta_t = R_{t+1} + \gamma V(S_{t+1}) - V(S_t)$  is called the *TD error*

## TD target for prediction

- The TD target:  $R_{t+1} + \gamma V(S_{t+1})$ 
  - it is an estimate like MC target because it samples the expected value
  - it is an estimate like the DP target because it uses the current estimate of  $V$




## Algorithm

```
Input: the policy  $\pi$  to be evaluated
Initialize  $V(s)$  arbitrarily (e.g.,  $V(s) = 0, \forall s \in \mathcal{S}^+$ )
Repeat (for each episode):
  Initialize  $S$ 
  Repeat (for each step of episode):
     $A \leftarrow$  action given by  $\pi$  for  $S$ 
    Take action  $A$ ; observe reward,  $R$ , and next state,  $S'$ 
     $V(S) \leftarrow V(S) + \alpha[R + \gamma V(S') - V(S)]$ 
     $S \leftarrow S'$ 
  until  $S$  is terminal
```

Figure 6.1: Tabular TD(0) for estimating  $v_\pi$ .

# Example



- 4 states MC: A, B, C, and a terminal state D. Assume the policy is random. The transition reward is -1,  $\alpha=0.5$ ,  $\gamma=1$ .

- Episode 1:  $A \rightarrow B \rightarrow C \rightarrow D$
- Episode 2:  $A \rightarrow C \rightarrow B \rightarrow C \rightarrow D$

- TD learning

- Episode 1:

- $V(A) = 0 + 0.5(-1 + 1 \times 0 - 0) = -0.5$
- $V(B) = 0 + 0.5(-1 + 1 \times 0 - 0) = -0.5$
- $V(C) = 0 + 0.5(-1 + 1 \times 0 - 0) = -0.5$
- $V(D) = 0$  (terminal)


- Episode 2:

- $V(A) = -0.5 + 0.5(-1 + 1 \times (-0.5) - (-0.5)) = -1.0$
- $V(C) = -0.5 + 0.5(-1 + 1 \times (-0.5) - (-0.5)) = -1.0$
- $V(B) = -0.5 + 0.5(-1 + 1 \times (-1.0) - (-0.5)) = -1.0$
- $V(C) = -1.0 + 0.5(-1 + 1 \times 0 - (-1.0)) = -1.5$

$$\text{Rule: } V(s) = V(s) + \alpha(R(s') + \gamma V(s') - V(s))$$


# Dynamic Programming Backup

$$V(S_t) \leftarrow \mathbb{E}_{\pi} [R_{t+1} + \gamma V(S_{t+1})]$$




# Monte-Carlo Backup

$$V(S_t) \leftarrow V(S_t) + \alpha (G_t - V(S_t))$$



# Temporal-Difference Backup

$$V(S_t) \leftarrow V(S_t) + \alpha (R_{t+1} + \gamma V(S_{t+1}) - V(S_t))$$



## Example: Driving Home

- Consider driving home:
  - Each day you drive home
  - Your goal is to try and predict how long it will take at particular stages
  - When you the leave office you note the time, day, & other relevant info
- Consider the policy evaluation or prediction task

## Example: Driving Home

<i>State</i>	<i>Elapsed Time (minutes)</i>	<i>Predicted Time to Go</i>	<i>Predicted Total Time</i>
leaving office, friday at 6	0	30	30
reach car, raining	5	35	40
exiting highway	20	15	35
2ndary road, behind truck	30	10	40
entering home street	40	3	43
arrive home	43	0	43

## Driving home as an RL problem


<i>State</i>	<i>Elapsed Time (minutes)</i>	<b>R</b>	<b>V(s)</b>	<b>V(office)</b>
			<i>Predicted Time to Go</i>	<i>Predicted Total Time</i>
leaving office, friday at 6	0	5	30	30
reach car, raining	5	15	35	40
exiting highway	20	10	15	35
2ndary road, behind truck	30	10	10	40
entering home street	40	3	3	43
arrive home	43		0	43

- **Goal:** update the prediction of total time leaving from office, while driving home
  - With **MC** we would need to **wait for a termination** – until we get home – then calculate Gt for each step of the episode, then apply our updates




# Driving home as an RL problem

Changes recommended by  
Monte Carlo methods ( $\alpha=1$ )



Changes recommended  
by TD methods ( $\alpha=1$ )



# Advantages and Disadvantages of MC vs. TD

## **TD can learn before knowing the final outcome**

- TD can learn online after every step
- MC must wait until the end of the episode before a return is known

## **TD can learn without the outcome**

- TD can learn from incomplete sequences
- MC can only learn from complete sequences
- TD works in continuing (non-terminating) environments
- MC only works for episodic (terminating) environments


**Both MC and TD converge (under certain assumptions), but which is faster?**

# Random Walk Example




- ▶ C is start state, episodic, undiscounted  $\gamma = 1$
- ▶  $\pi$  is left or right with equal probability in all states
- ▶ termination at either end
- ▶ rewards +1 on **right** termination, 0 otherwise
- ▶ what does  $v_{\pi}(s)$  tell us?
  - probability of termination on right side from each state, under random policy
  - what is  $v_{\pi} = [A \ B \ C \ D \ E]$ ?
    - $v_{\pi} = [1/6 \ 2/6 \ 3/6 \ 4/6 \ 5/6]$
- ▶ Initialize  $V(s) = 0.5 \ \forall \ s \in \mathcal{S}$

## TD and MC



$$V(S_t) \leftarrow V(S_t) + \alpha [R_{t+1} + \gamma V(S_{t+1}) - V(S_t)]$$




## Batch MC and TD

- **Batch Updating:** train completely on a finite amount of data
  - e.g., train repeatedly on 10 episodes until convergence.
- Compute updates according to TD or MC, but only update estimates after each complete pass through the data.
- For any finite Markov prediction task, under batch updating, TD converges for sufficiently small  $\alpha$ .
- Constant- $\alpha$  MC also converges under these conditions, but to a different answer!

## Random Walk under Batch Updating

- After each new episode, all episodes seen so far are treated as a batch
- This growing batch is repeatedly processed by TD and MC until convergence



End of lecture